



Ensō Operating System

The Kernel

Scheduler: Round-Robin, Priority Queues, Tree Flavours

Scheduler Actors: Features, Timers, Async I/O

Streams Backends: Zero-copy, Message Passing

Linear Backends: Async I/O Disk Streams, Network Streams

Indexed Backends: Timers, Actors

Backpressured Message Bus/Buffers: Arc/Vec prealloc

Class: Low Latency, Real Time

Respect Kernels History

Richard Rashid. Mach 3. NUMA, Bus Oriented Components
Dave Cutler. Windows NT. True Async I/O on IoCompletionPort
BeOS. Travis Geiselbrecht. SMP Scaling of OS services
Microkernels: RT, Tiny codebases eCos/TRON, QNX, VxWorks
Unikernels: Erlang, Mirage, HaLVM

FOUNDATION

Stream/List duality

```
pub enum List<Message> {  
    Nil,  
    Cons(Message, std::marker::PhantomData<List<Message>>) }
```

```
pub struct Stream<Message> {  
    head: Message,  
    tail: Box<Stream<Message>> }
```

```
pub trait Stream<Message> {  
    fn head(&mut self) -> Message;  
    fn tail(&mut self) -> Stream<Message>; }
```

FUTURES

Zero-copy and Message Passing

```
pub trait Future<Message,Error> {  
    fn poll(&mut self) -> Result<Message,Error>;  
    fn tail(&mut self) -> Future<Message,Error>; }  

```

```
pub trait Process<Protocol, State, Error> {  
    fn state(&mut self) -> State;  
    fn send(&mut self, Protocol) -> Result<State, Error>; }  

```

```
pub trait Discipline<Stream<Message>> {  
    fn select(&mut self, u64) -> Stream<Message>; }  

```

ORDERS

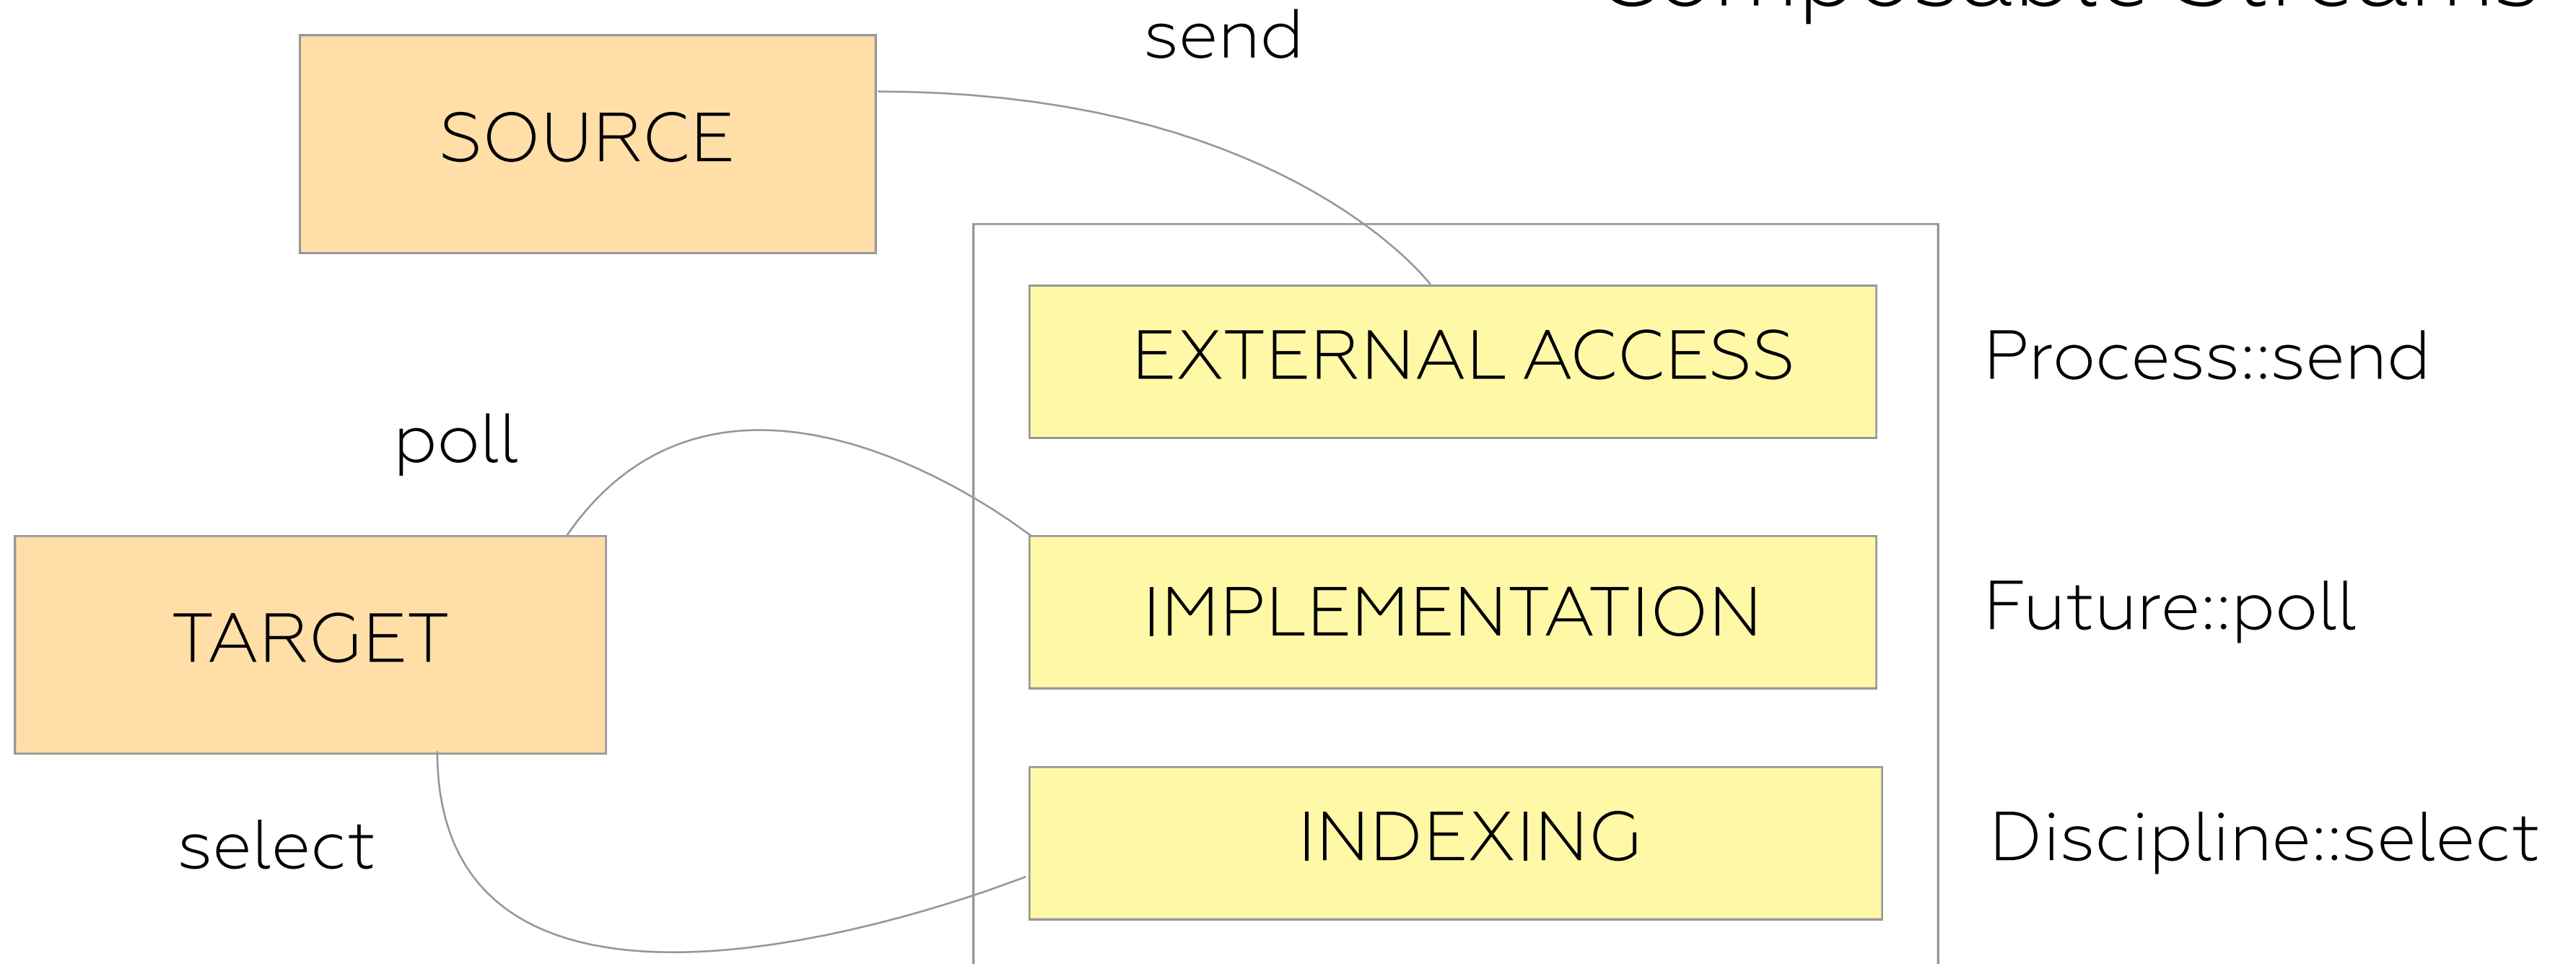
Order Processing Protocol

```
pub struct OrderState {  
    state: OrderStatus,  
    id: ID,  
    price: Price,  
    size: Size,  
    side: Side, }
```

```
pub enum OrderProtocol {  
    Request, Execute, Reject,  
    UnsolicitedCancel,  
    Cancel,  
    Ack, Replace, }
```

```
type OrderProcess = Process<OrderProtocol, OrderState, Error>;
```

Composable Streams



Polymorphic Disciplines

SCHED


DISK

Discipline::select

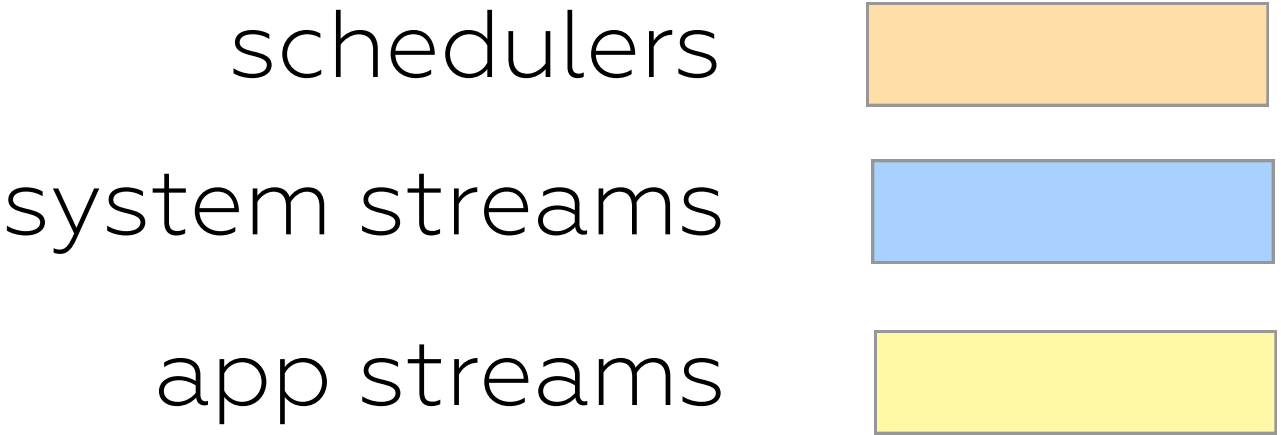
Circular Buffers

TIMERS

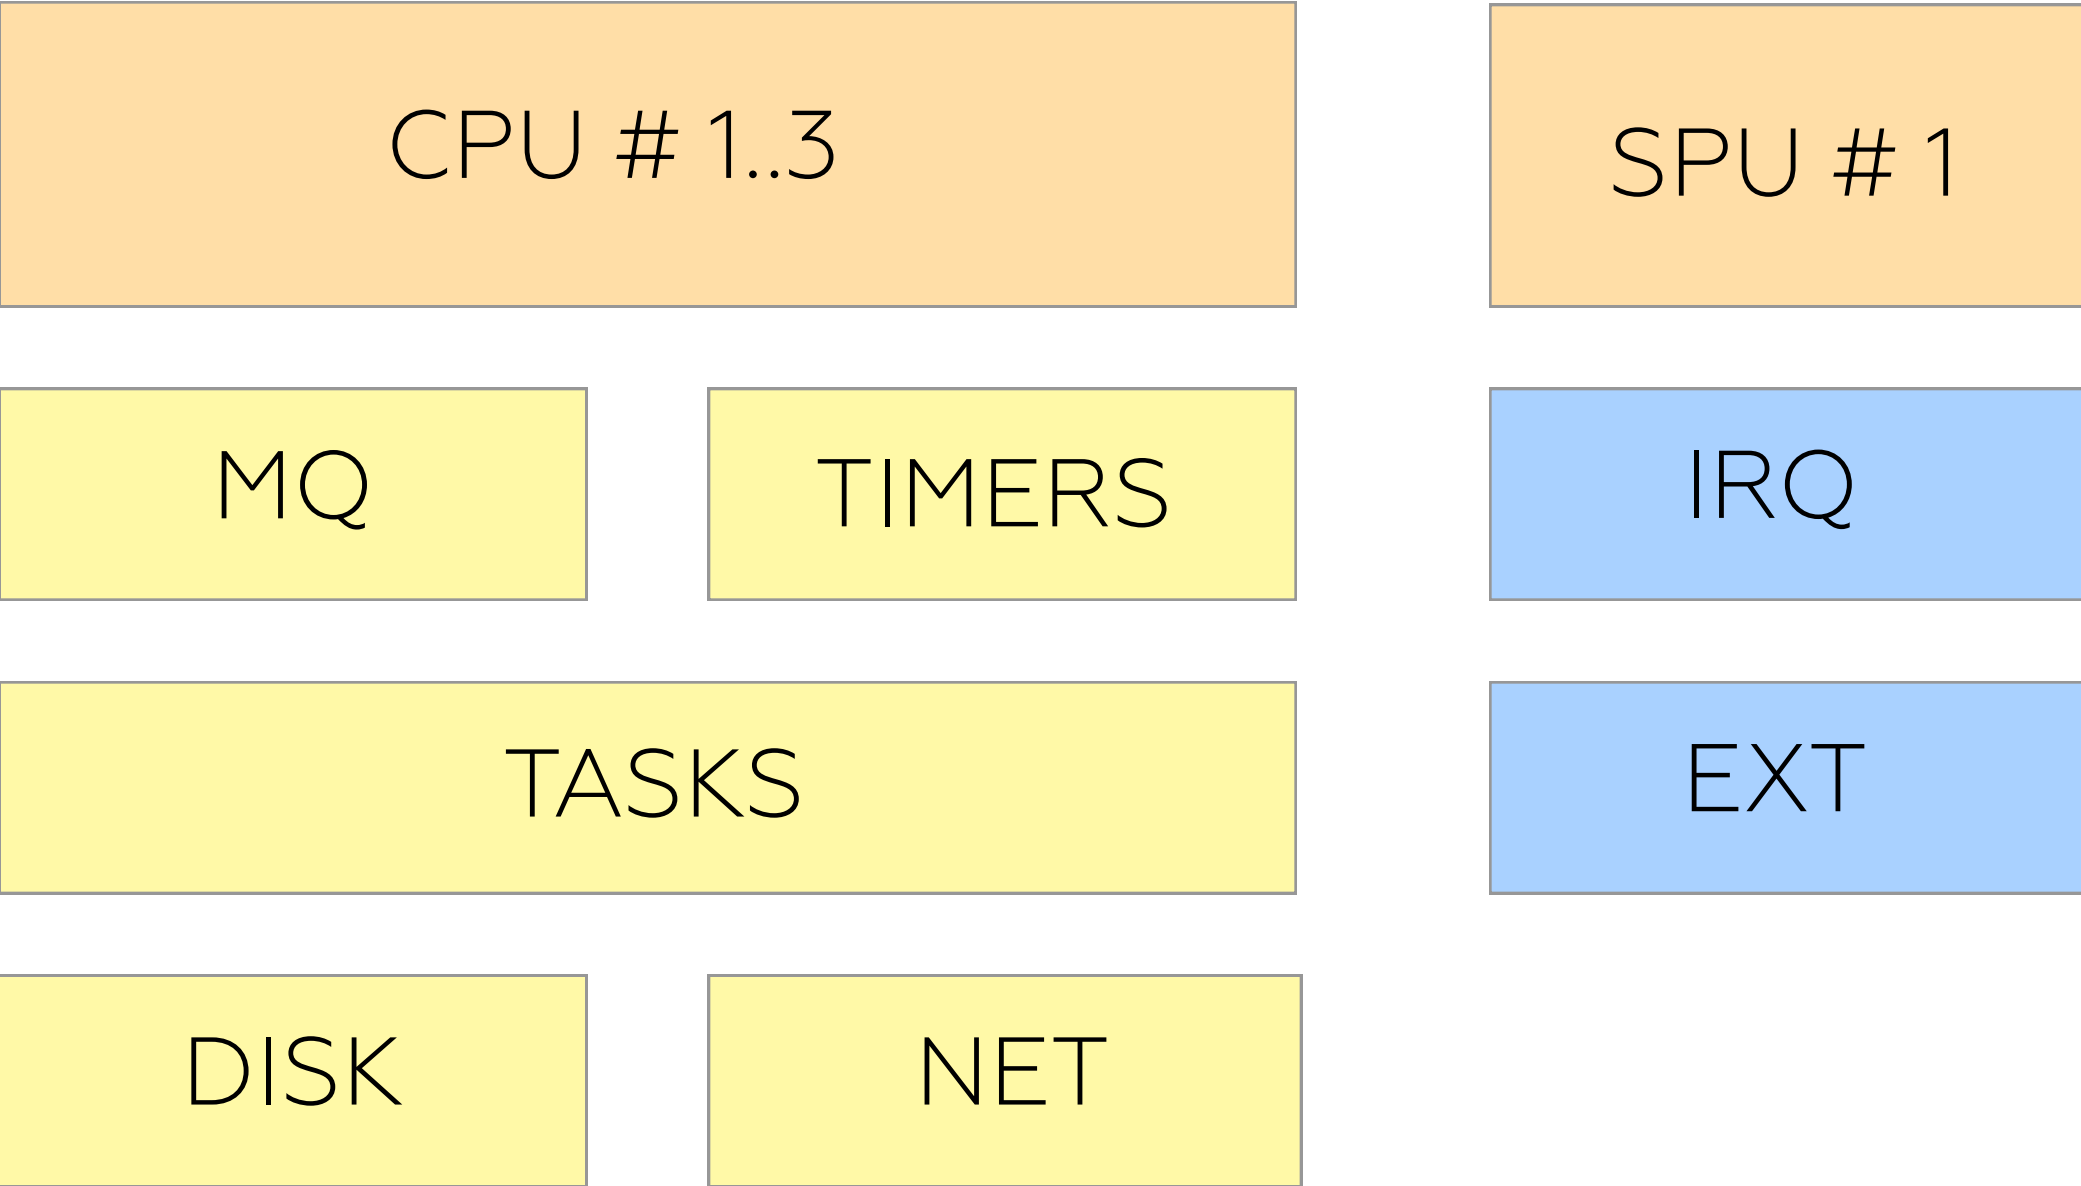
Discipline::select



Linear: MQ, EXT, DISK, NET
Trees: TIMERS
Priority Queues: TASKS, IRQ



Node Components



REACTOR

Core Context

TOKIO

REVENT

NANOMSG

RX/RUST

Discipline::select

TASKS

Single Task Tree per Core

Scheduler::select

Reactor Requested Features

1. Reactor Core Context
2. Task Context
3. Scheduleable Entities
4. Task List
5. Task Selection by Priority
6. API: spawn/select/poll
7. Inspirational Libraries: Rx, Nanomsg, Tokio

MQ

Queue Types

SPSC/LINK

MPSC/SUB

SPMC/PUB

Vec, VecDeque, Link, Turbine, Pipes

(rx , tx)

CHANNEL

Discipline::select

Non-schedulable Array Store

Queues Requested Features

1. Three Basic Drivers: SPSC, MPSC, SPSC
2. Different Backends: Vec, VecDeque, Link, Turbine, Pipes
3. Queue Properties
4. Queue with Priorities
5. API: create/pub/sub/link
6. Inspirational Libraries: Turbine, Rust Queues, Pipes

TIMERS

Discipline::select

TIMER

(interval , task_id)

TOKIO

Dedicated Core or Interrupt Queue

Scheduler::select

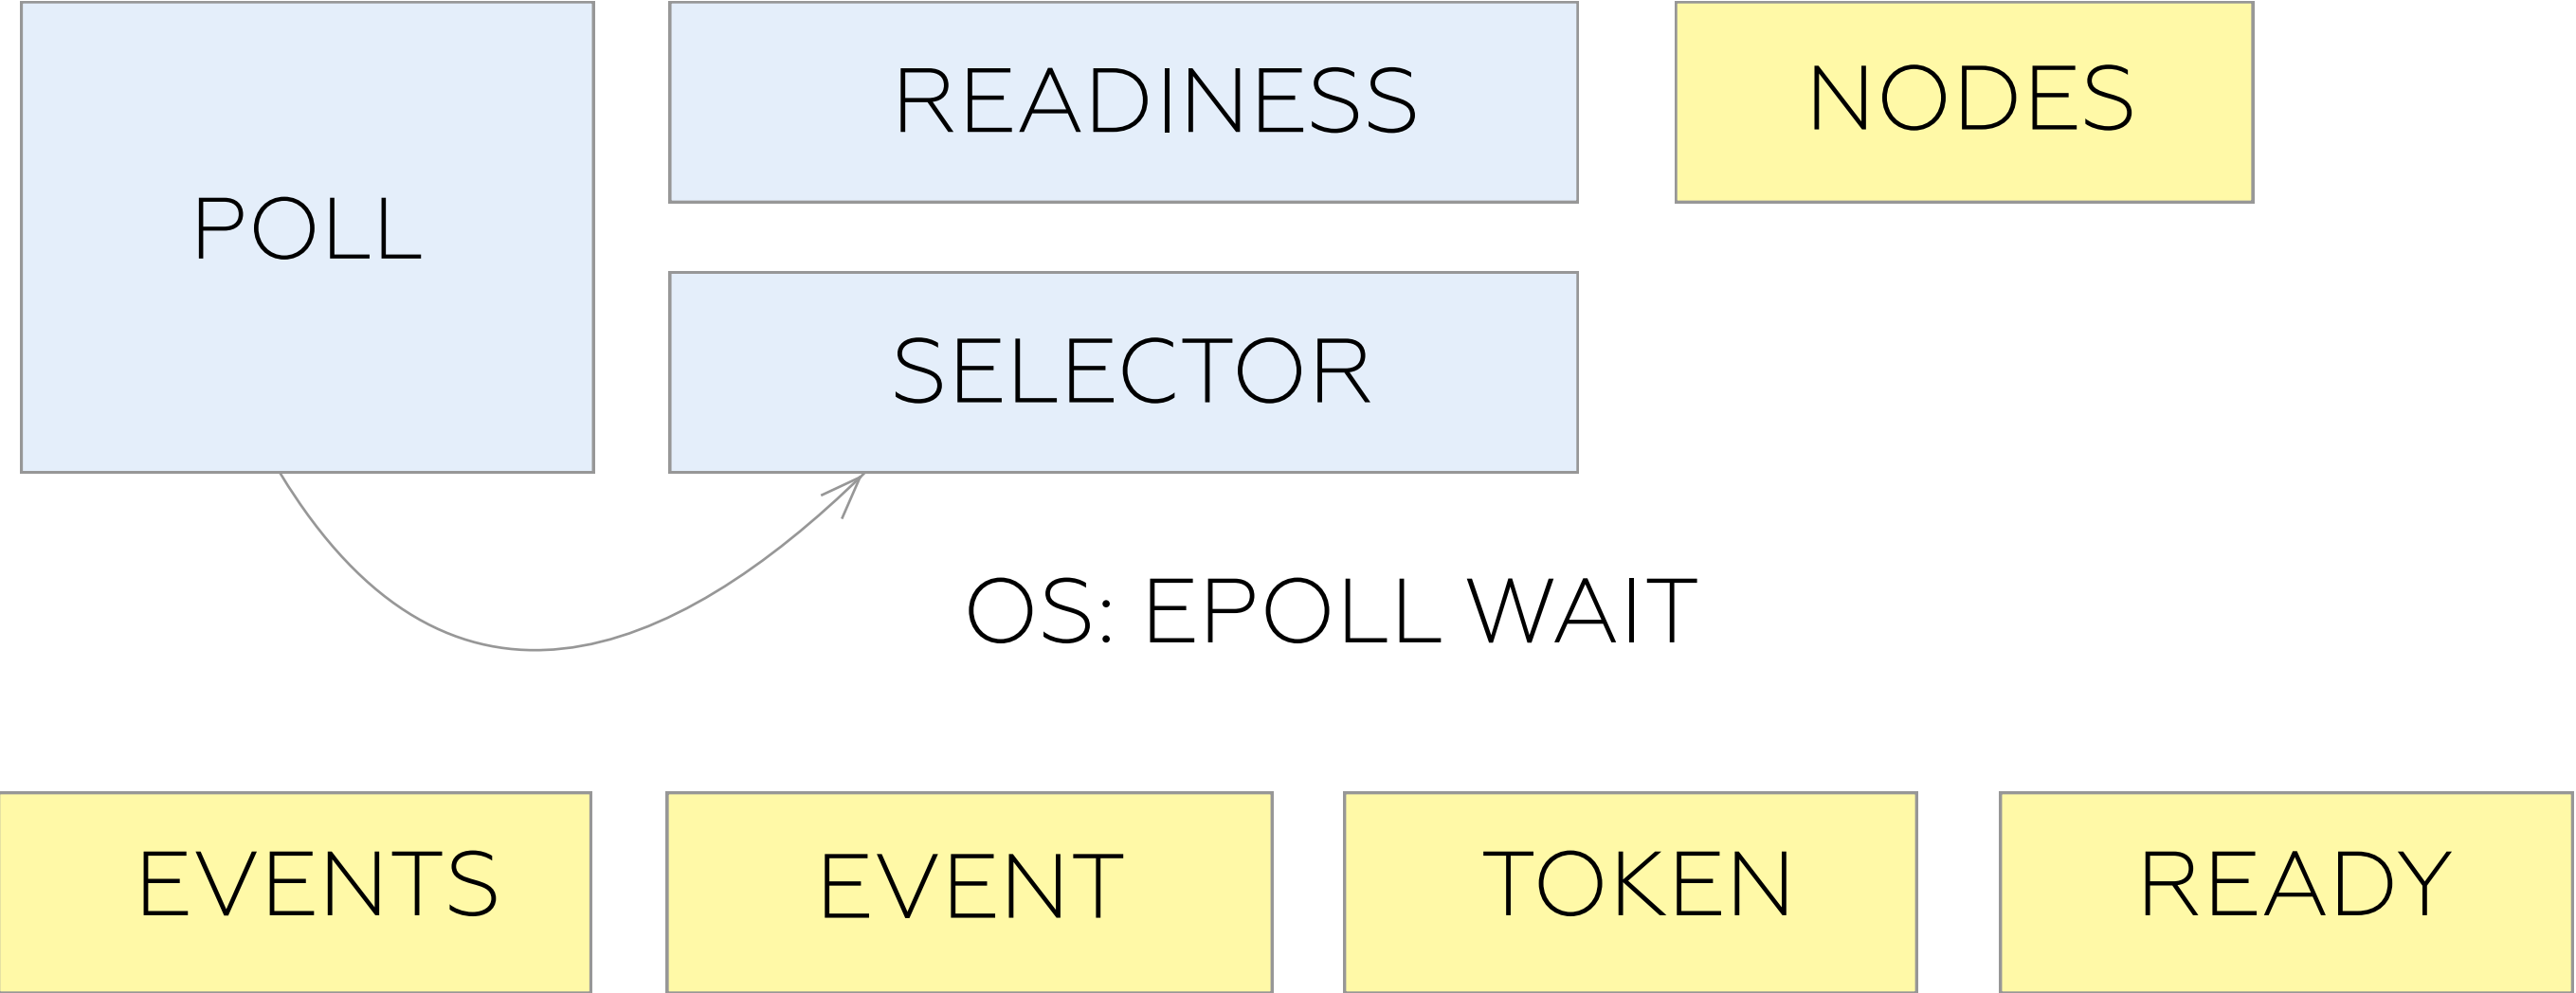
RX/RUST

Timers Requested Features

1. Timers Core Context
2. Timer Types: Oneshot Scalable Timers, Fixed List of Interval Timers
3. API: create/select/poll
4. Inspirational Libraries: Tokio, Rx

NET I/O

MIO compatible polling loop based on Readiness Queue

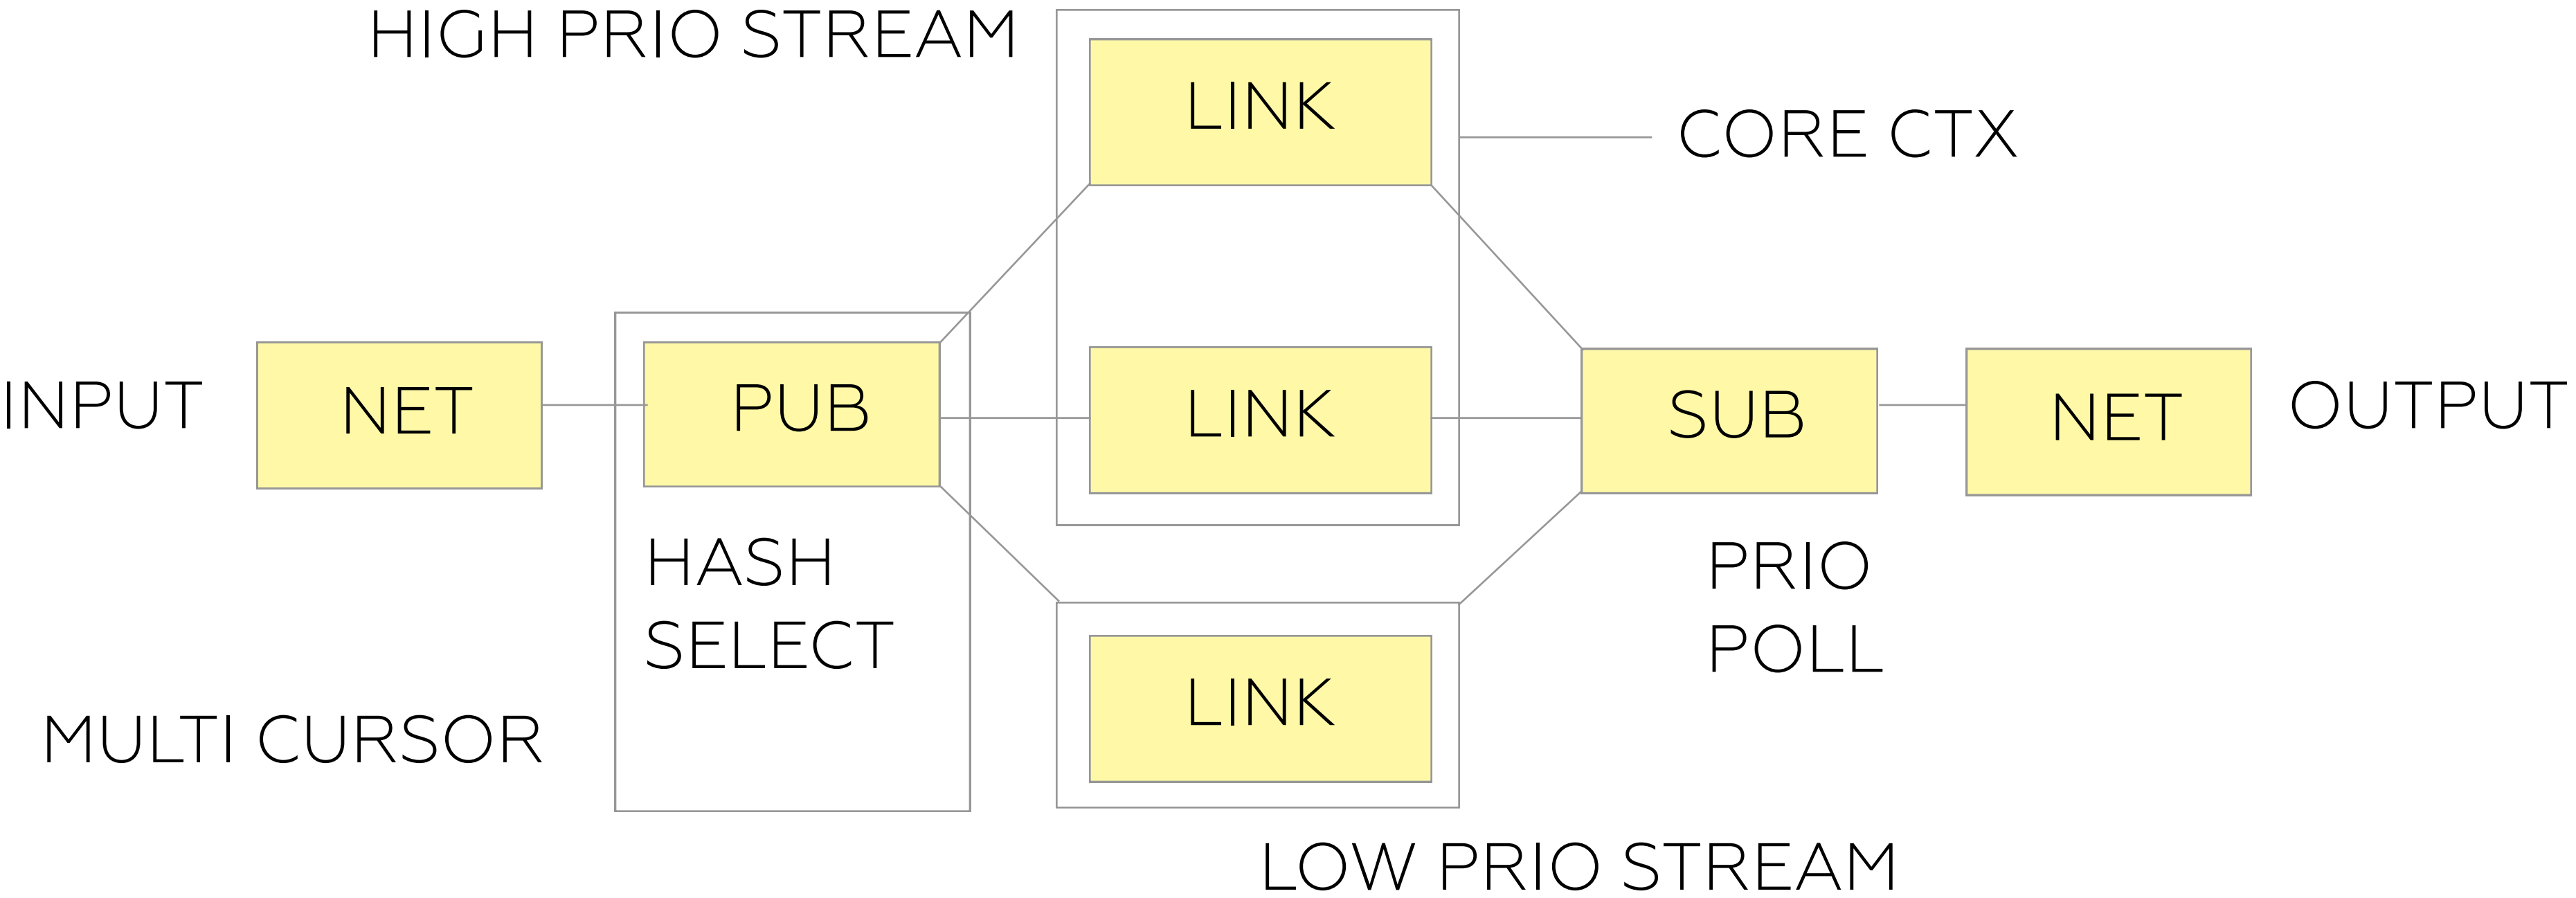


Network I/O Features

1. MIO compatible API
2. Evented Polling on Linux, BSD, XSOCK
3. Connection State
4. API: create/select/poll
5. Close to FFI API
6. Client and Server Samples

LOAD BALANCING CASE

Load Balancing
of Priority Streams per Core Buckets



FAST DELIVERY CASE

Single Threaded Task Configuration
to be compared as reference

I/O TASK



CPU TASK



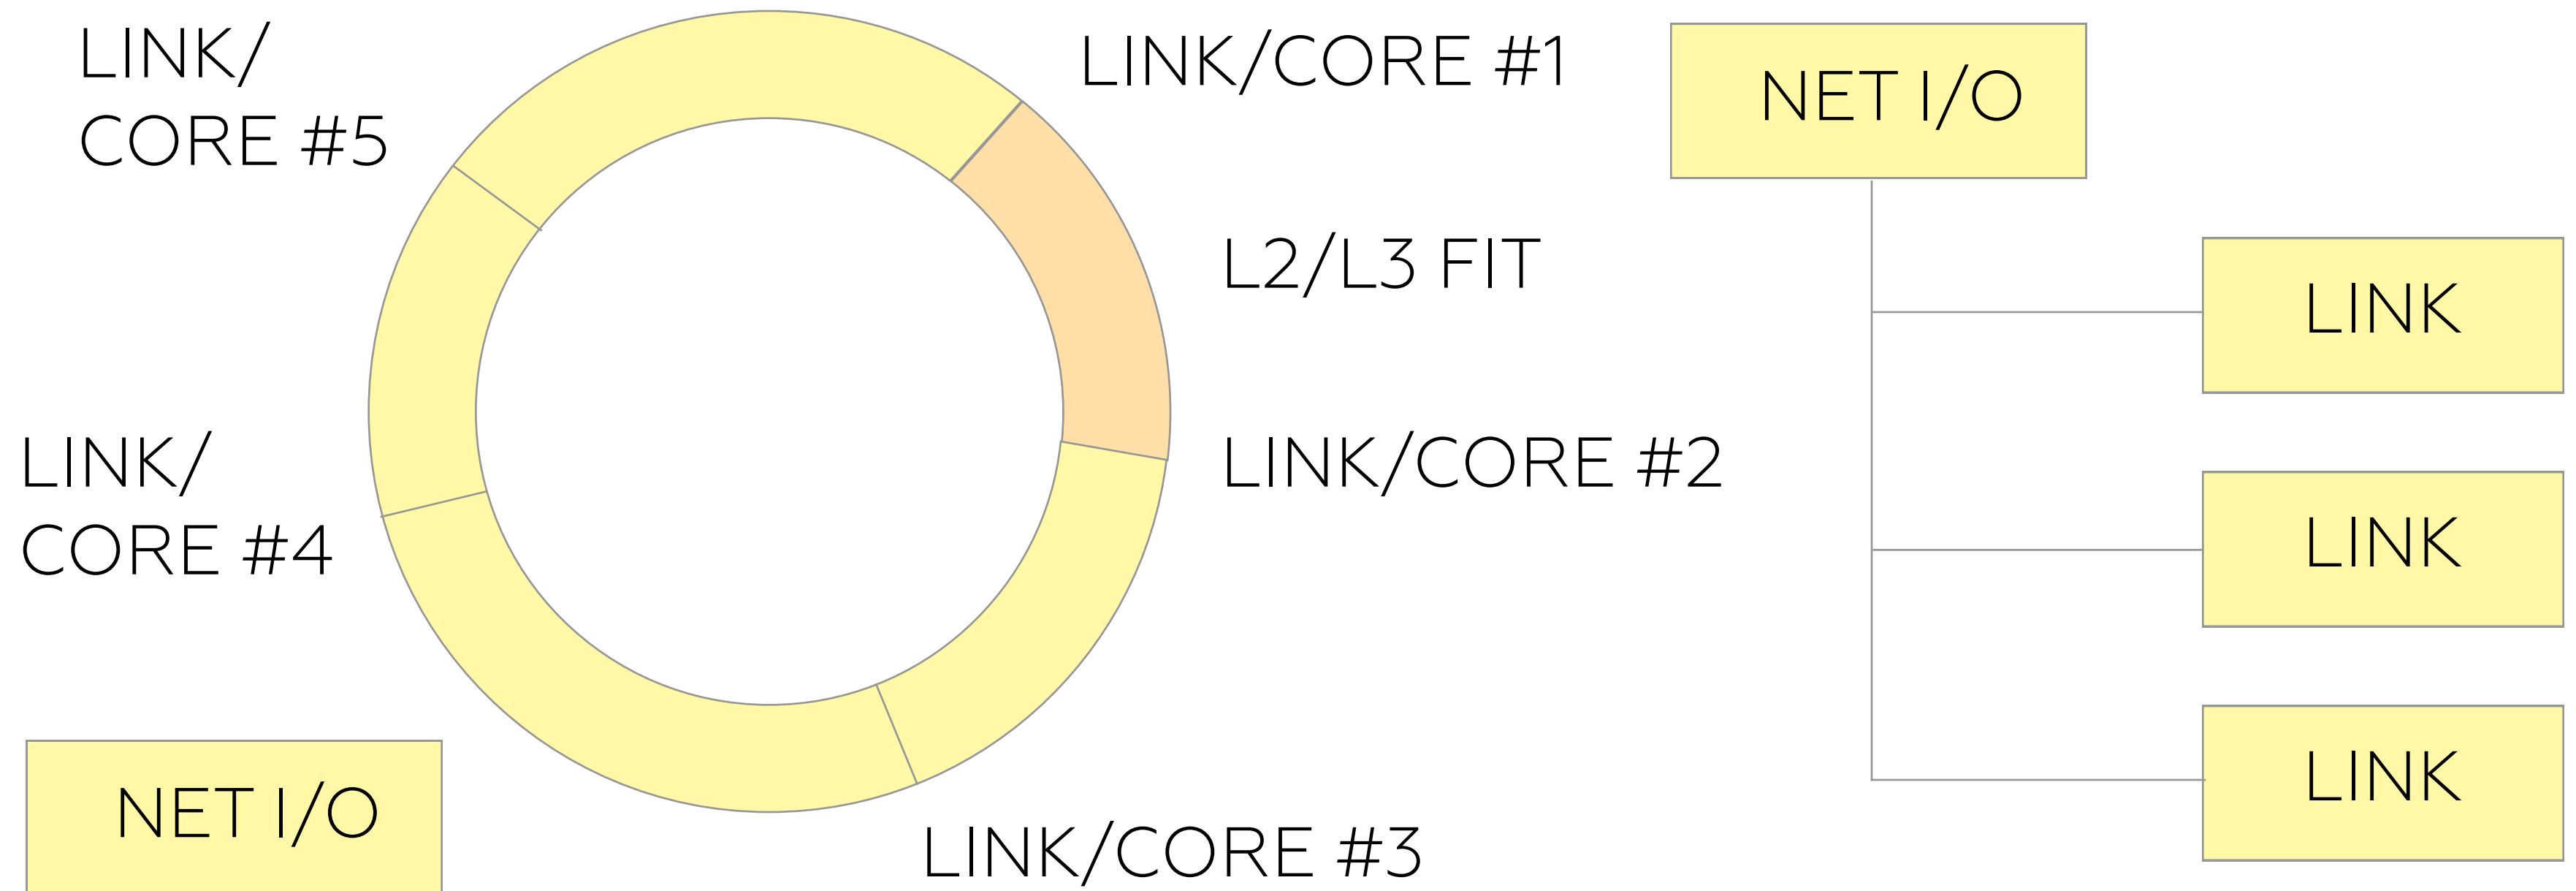
I/O TASK



You can use inplace message modifying and reduce copies to unpack and pack.

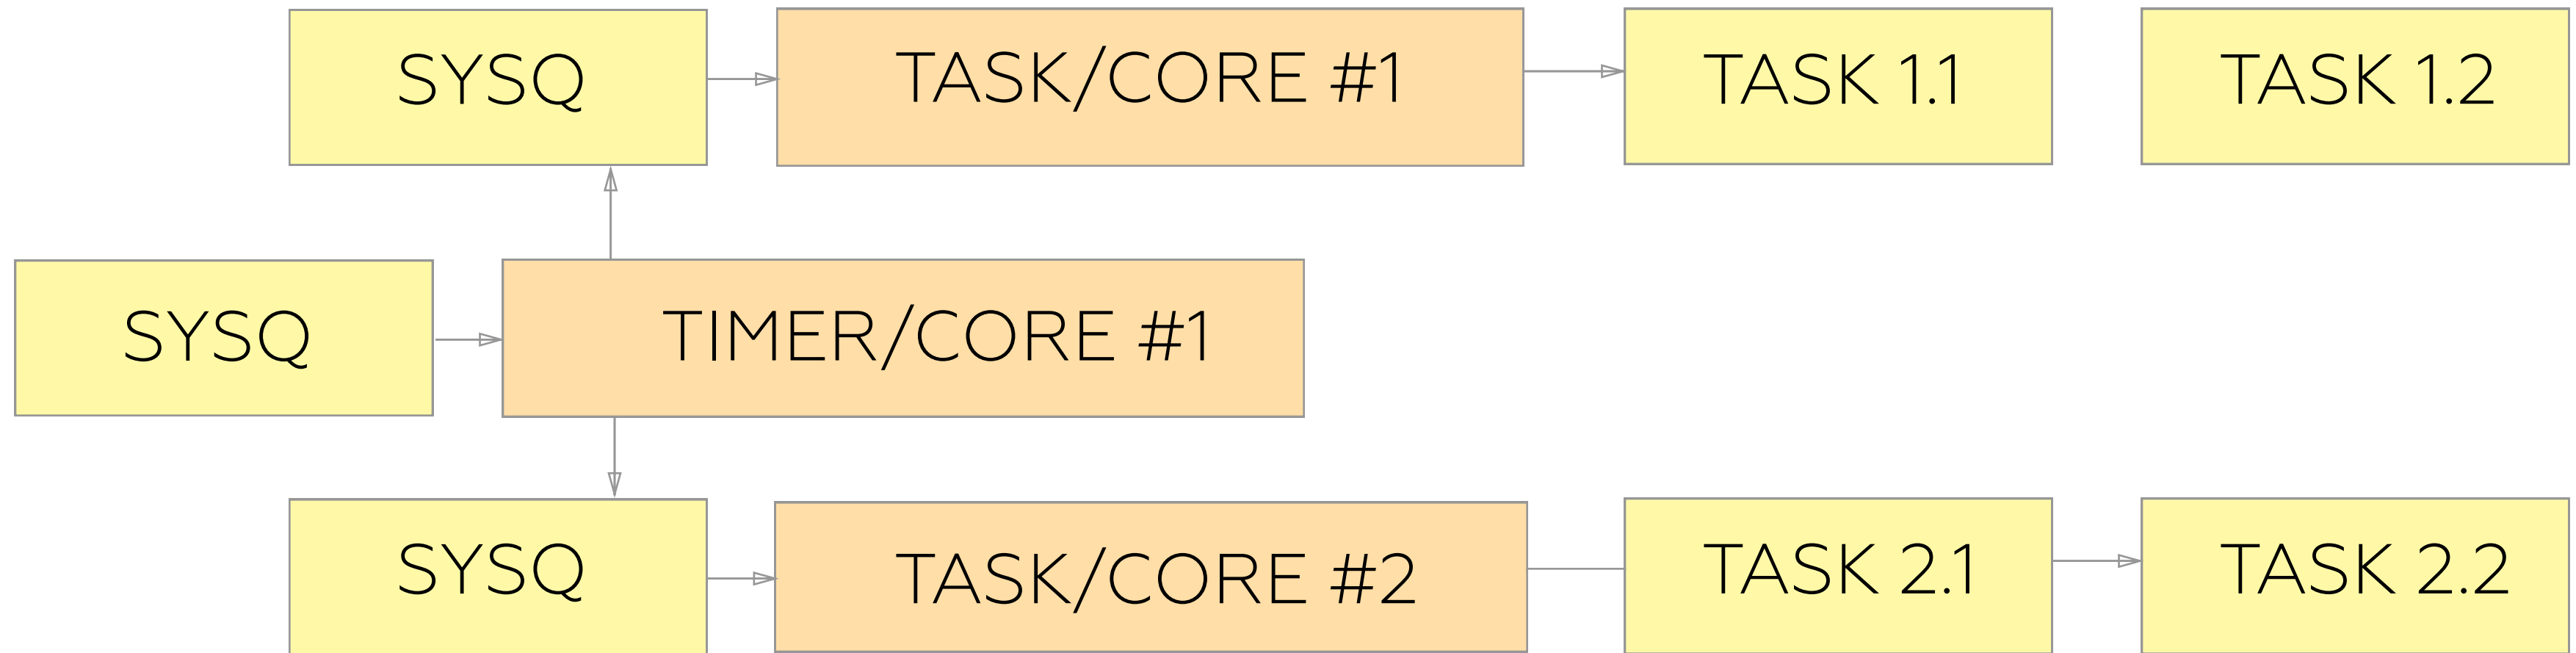
PUBLISHER CASE

Multicursor Implementation of PUB
for InterCore Queue Migrations and Cache Locality



INTERCORE QUEUES

Scheduler Reactors can communicate through InterCore transport when needed.



Network Polling also could be seen as InterCore single sourced Queue for Task Delivery.