

The Kernel

Scheduler: Round-Robin, Priority Queues, Tree Flavours

Scheduler Actors: Features, Timers, Async I/O

Streams Backends: Zero-copy, Message Passing

Linear Backends: Async I/O Disk Streams, Network Streams

Indexed Backends: Timers, Actors

Backpressured Message Bus/Buffers: Arc/Vec prealloc

Class: Low Latency, Real Time

Respect Kernels History

Richard Rashid. Mach 3. NUMA, Bus Oriented Components
Dave Cutler. Windows NT. True Async I/O on IoCompletionPort
BeOS. Travis Geiselbrecht. SMP Scaling of OS services
Microkernels: RT, Tiny codebases eCos/TRON, QNX, VxWorks
Unikernels: Erlang, Mirage, HaLVM

FOUNDATION

Stream/List duality

```
pub enum List<Message> {  
    Nil,  
    Cons(Message, std::marker::PhantomData<List<Message>>) }
```

```
pub struct Stream<Message> {  
    head: Message,  
    tail: Box<Stream<Message>> }
```

```
pub trait Stream<Message> {  
    fn head(&mut self) -> Message;  
    fn tail(&mut self) -> Stream<Message>; }
```

FUTURES

Zero-copy and Message Passing

```
pub trait Future<Message,Error> {  
    fn poll(&mut self) -> Result<Message,Error>;  
    fn tail(&mut self) -> Future<Message,Error>; }
```

```
pub trait Process<Protocol, State, Error> {  
    fn state(&mut self) -> State;  
    fn send(&mut self, Protocol) -> Result<State, Error>; }
```

```
pub trait Discipline<Stream<Message>> {  
    fn select(&mut self, u64) -> Stream<Message>; }
```

ORDERS

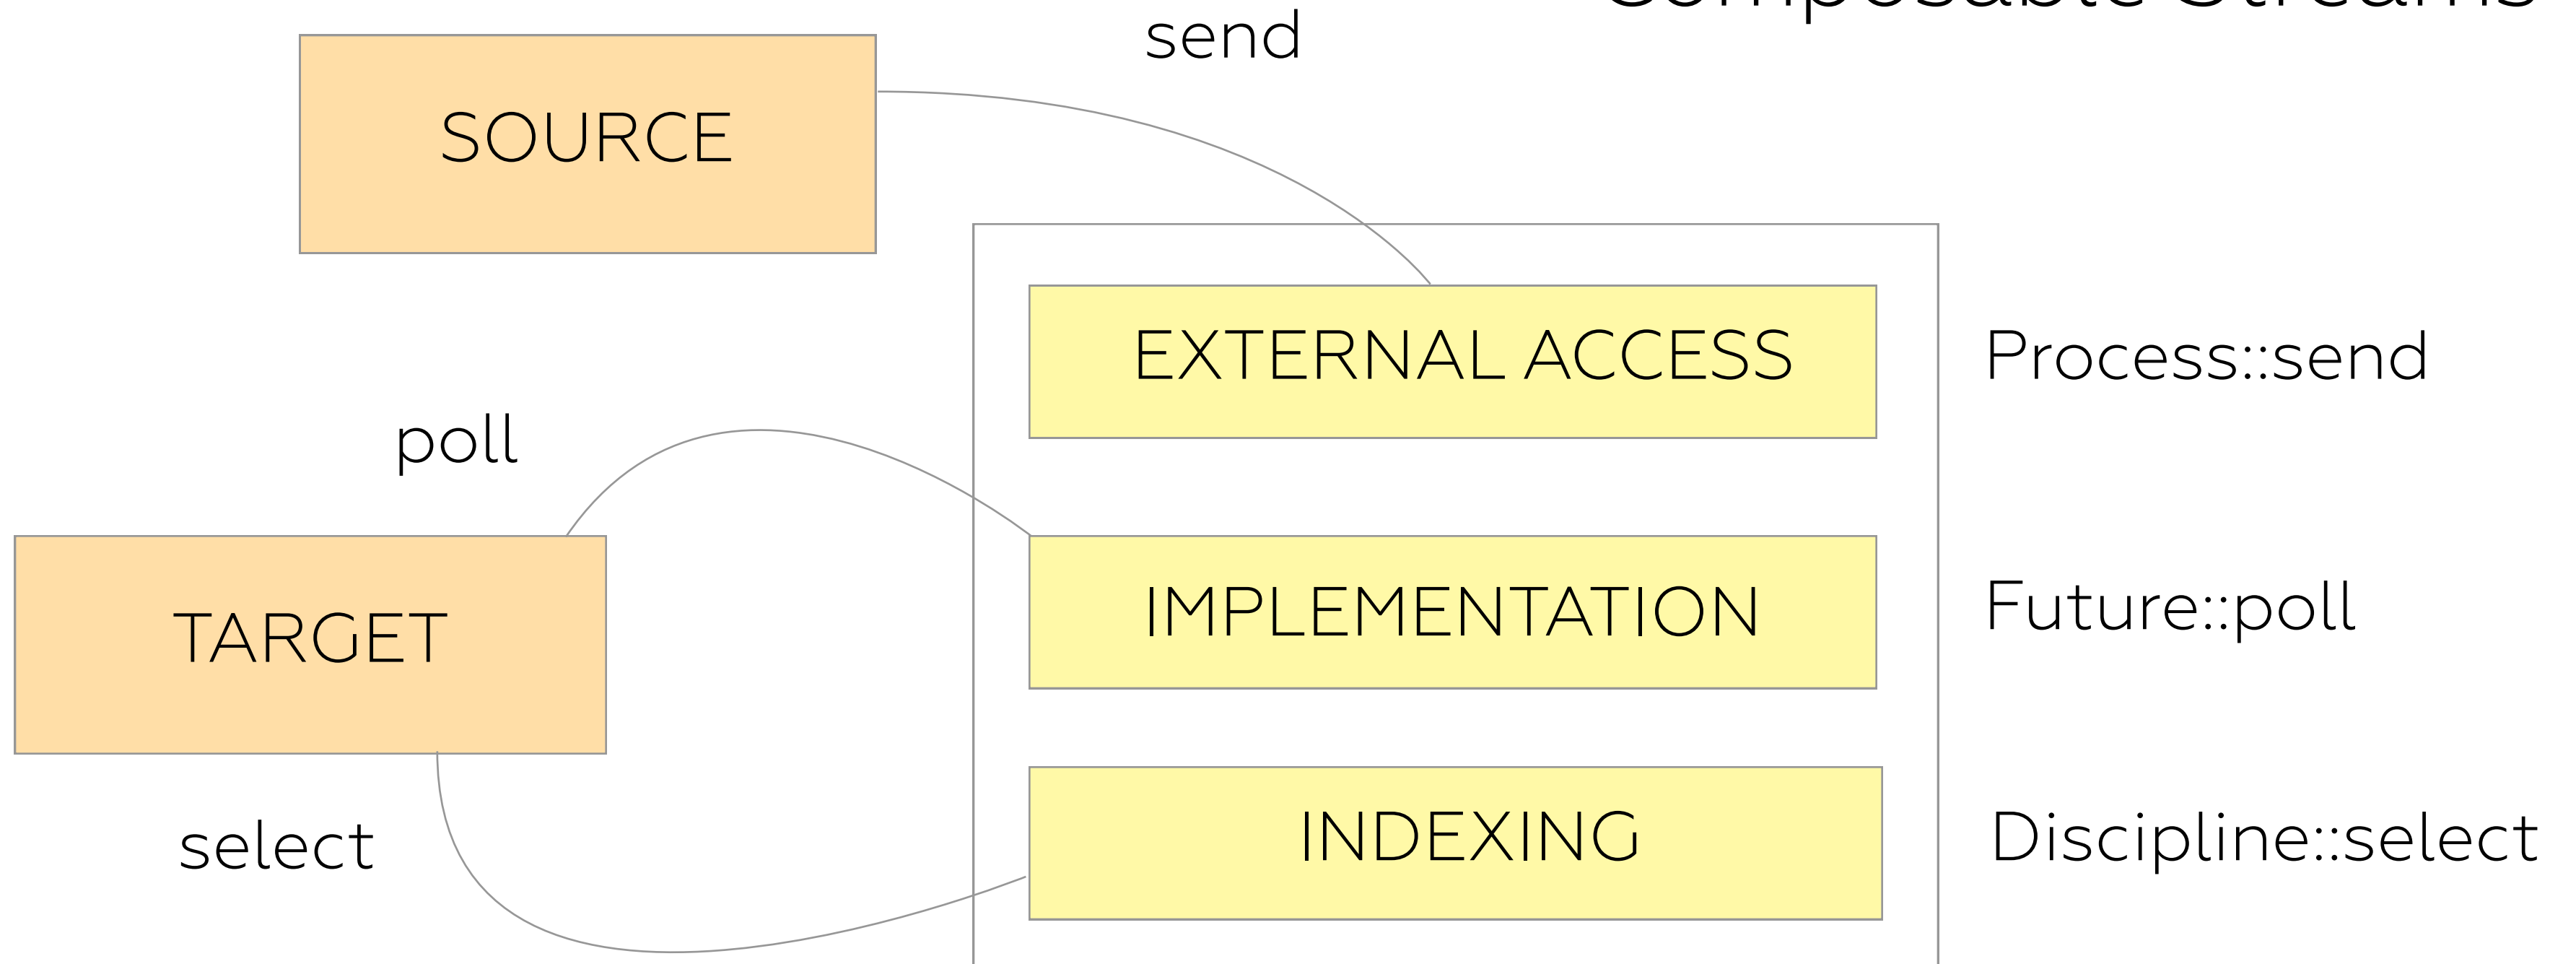
Order Processing Protocol

```
pub struct OrderState {  
    state: OrderStatus,  
    id: ID,  
    price: Price,  
    size: Size,  
    side: Side, }
```

```
pub enum OrderProtocol {  
    Request, Execute, Reject,  
    UnsolicitedCancel,  
    Cancel,  
    Ack, Replace, }
```

```
type OrderProcess = Process<OrderProtocol, OrderState, Error>;
```

Composable Streams



Polymorphic Disciplines

SCHED

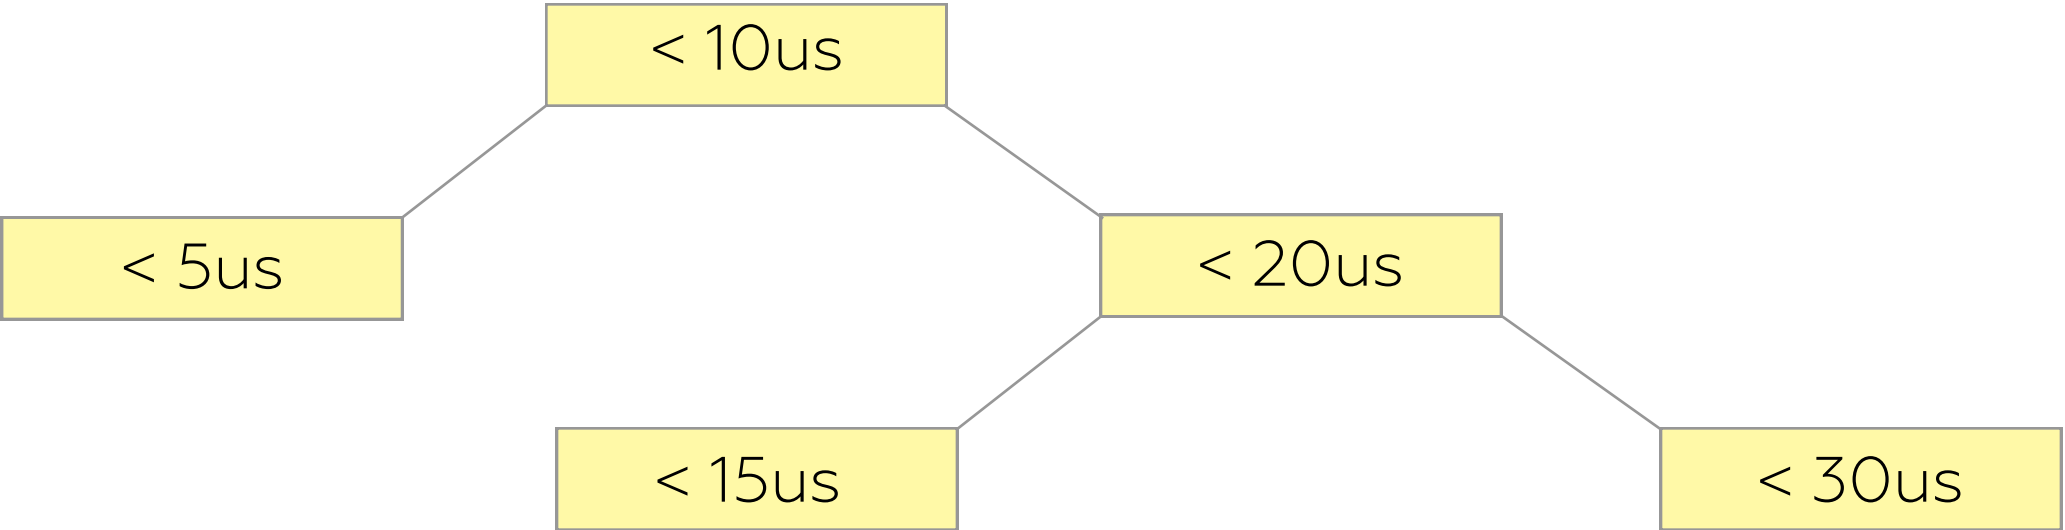
DISK

Discipline::select

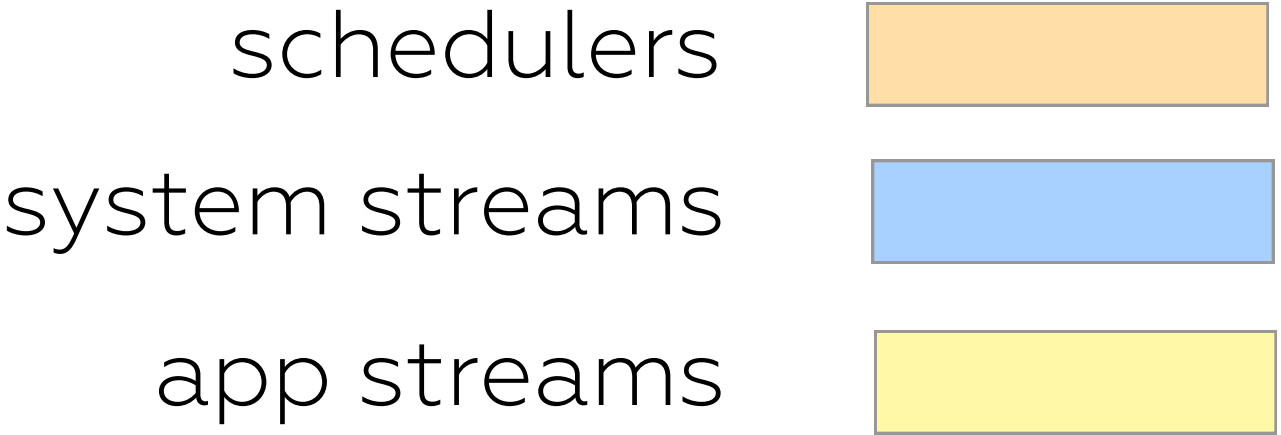
Circular Buffers

TIMERS

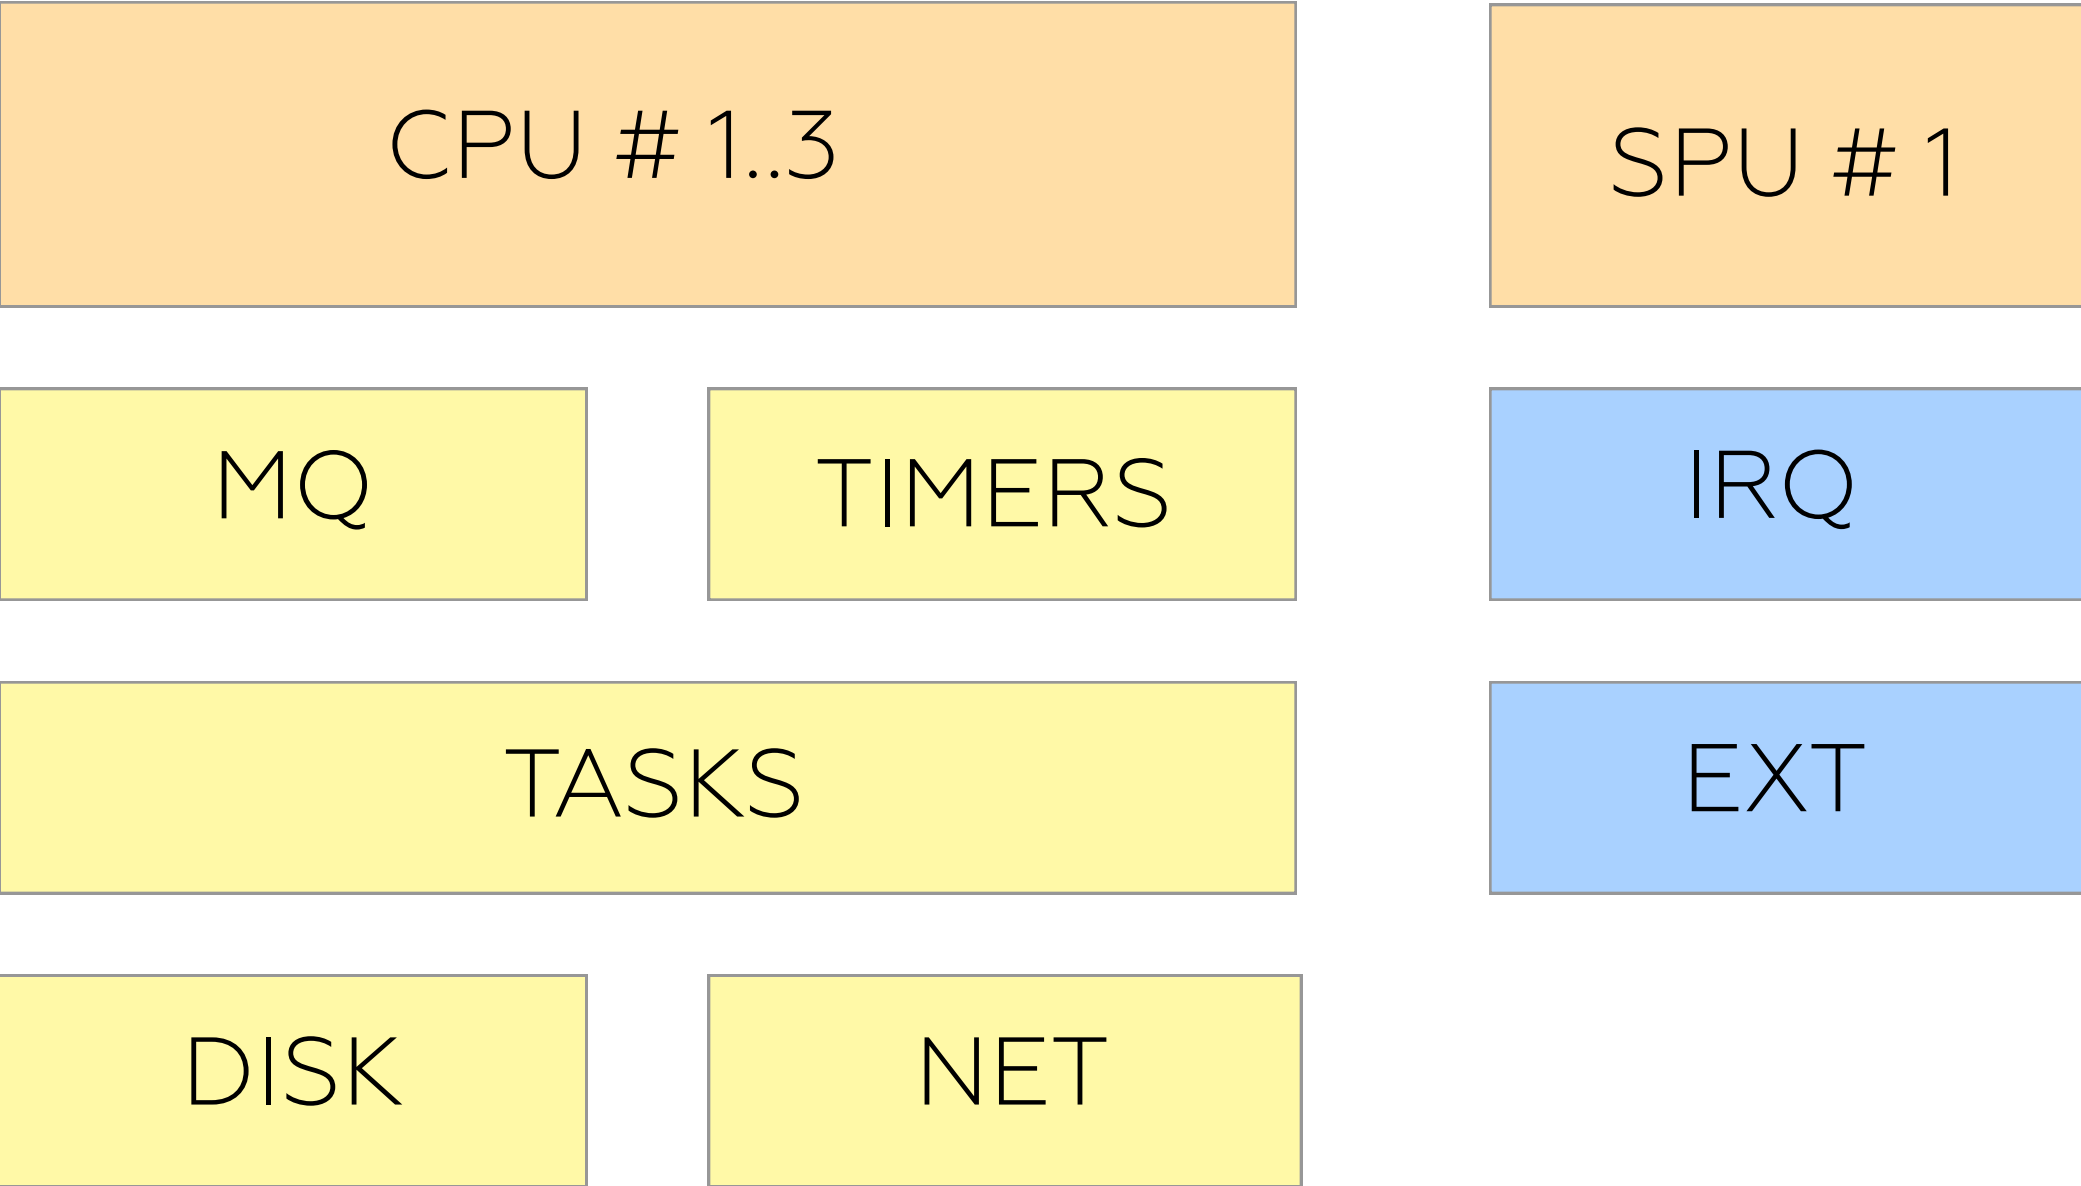
Discipline::select



Linear: MQ, EXT, DISK, NET
Trees: TIMERS
Priority Queues: TASKS, IRQ



Node Components



REACTOR

Discipline::select

TASKS

Core Context

TOKIO

NANOMSG

RX/RUST

Single Task Tree per Core

Scheduler::select

Reactor Requested Features

1. Reactor Core Context
2. Task Context
3. Scheduleable Entities
4. Task List
5. Task Selection by Priority
6. API: spawn/select/poll
7. Inspirational Libraries: Rx, Nanomsg, Tokio

MQ

Queue Types

SPSC/LINK

MPSC/SUB

SPMC/PUB

Vec, VecDeque, Link, Turbine, Pipes

CHANNEL

Discipline::select

Non-schedulable Array Store

Queues Requested Features

1. Three Basic Drivers: SPSC, MPSC, SPSC
2. Different Backends: Vec, VecDeque, Link, Turbine, Pipes
3. Queue Properties
4. Queue with Priorities
5. API: create/pub/sub/link
6. Inspirational Libraries: Turbine, Rust Queues, Pipes

TIMERS

Discipline::select

TIMER

(interval , task_id)

TOKIO

Dedicated Core or Interrupt Queue

Scheduler::select

RX/RUST

Timers Requested Features

1. Timers Core Context
2. Timer Types: Oneshot Scalable Timers, Fixed List of Interval Timers
3. API: create/select/poll
4. Inspirational Libraries: Tokio, Rx

NET I/O

Socket Context

TOKIO

NANOMSG

Vec

BUFFER

Single Chain per
Core

Scheduler::select

Network I/O Requested Features

1. Network I/O Core Context
2. Socket Context
3. Pipes/Queues Context and Interoperability
4. Timer Types: Oneshot Scalable Timers, Fixed List of Interval Timers
5. API: create/select/poll
6. Reconnections
7. Inspirational Libraries: Tokio, Rx